





## Map Reduce 2.0 Developing First MapReduce Job

Originals of Slides and Source Code for Examples:

<http://www.coreservlets.com/hadoop-tutorial/>

**Customized Java EE Training:** <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.  
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.



**For live Hadoop training, please see courses  
at <http://courses.coreservlets.com/>.**

**Taught by the author of this Hadoop tutorial. Available  
at public venues, or customized versions can be held  
on-site at your organization.**

- Courses developed and taught by Marty Hall
    - JSF 2, PrimeFaces, servlets/JSP, Ajax, jQuery, Android development, Java 6 or 7 programming, custom mix of topics
    - Ajax courses can concentrate on 1 library (jQuery, Prototype/Scriptaculous, Ext-JS, Dojo, etc.) or survey several
  - Courses developed and taught by [coreservlets.com](http://coreservlets.com) experts (edited by Marty)
    - **Hadoop**, Spring, Hibernate/JPA, GWT, SOAP-based and RESTful Web Services
- Contact [hall@coreservlets.com](mailto:hall@coreservlets.com) for details**

# Agenda

- **Introduce MapReduce framework**
- **Implement first MapReduce Job**

4

# MapReduce

- **Divided in two phases**
  - Map phase
  - Reduce phase
- **Both phases use key-value pairs as input and output**
- **The implementer provides map and reduce functions**
- **MapReduce framework orchestrates splitting, and distributing of Map and Reduce phases**
  - Most of the pieces can be easily overridden

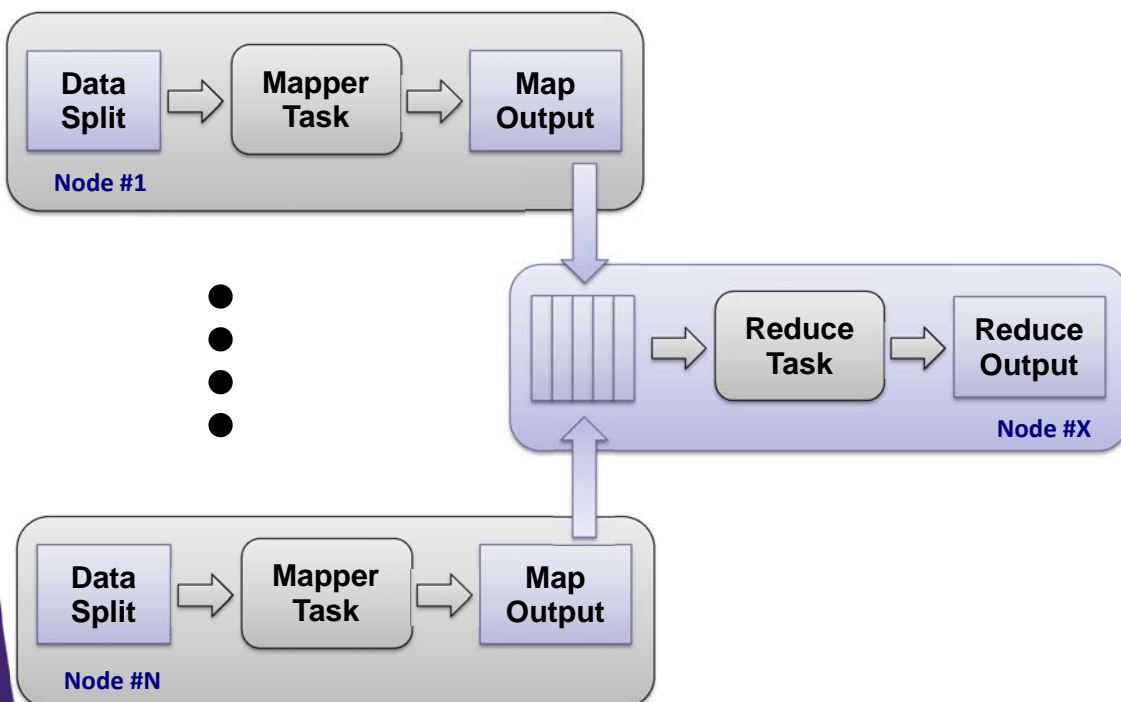
5

# MapReduce

- **Job** – execution of map and reduce functions to accomplish a task
  - Equal to Java's main
- **Task** – single Mapper or Reducer
  - Performs work on a fragment of data

6

## Map Reduce Flow of Data



7

# First Map Reduce Job

- **StartsWithCount Job**

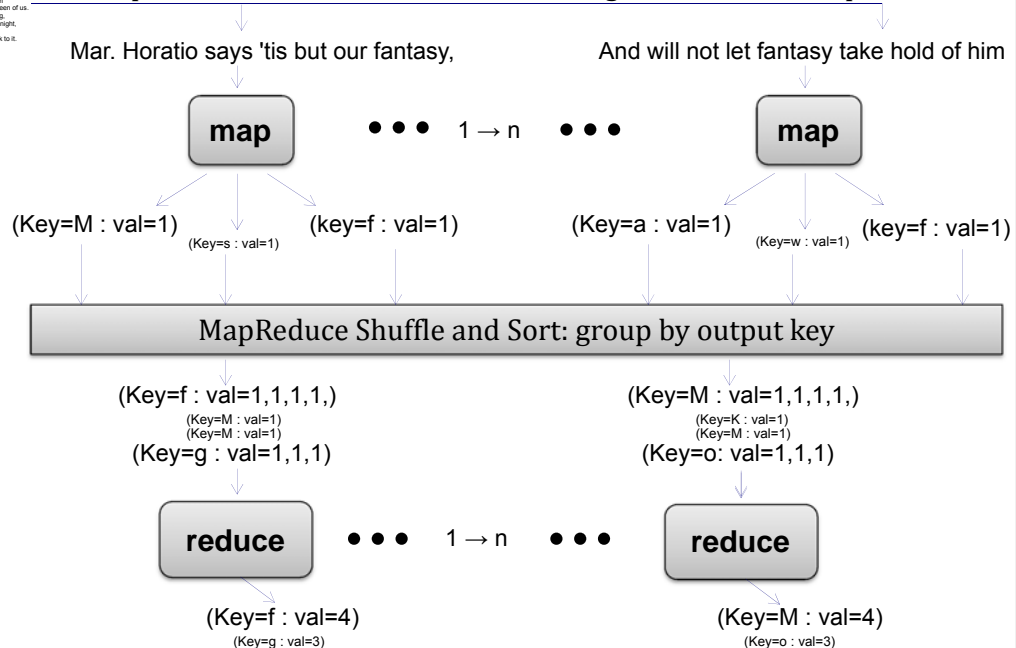
- Input is a body of text from HDFS
  - In this case hamlet.txt
- Split text into tokens
- For each first letter sum up all occurrences
- Output to HDFS

8

# Word Count Job

Mar. What! has this thing appear'd again to-night?  
Hor. I have seen nothing.  
Mar. Horatio says 'tis but our fantasy,  
And will not let our better parts of him  
Touching this deadèd sight, twice seen of us.  
Therefore I have entreated him along;  
With us to watch the minutes of this night;  
That, if again this apparition come,  
He may approve our eyes and speak to it.  
Hor. Tush, tush, 'twill not appear.

MapReduce breaks text into lines feeding each line into map functions



9

# StartsWithCount Job

## 1. Configure the Job

- Specify Input, Output, Mapper, Reducer and Combiner

## 2. Implement Mapper

- Input is text – a line from hamlet.txt
- Tokenize the text and emit first character with a count of 1 - <token, 1>

## 3. Implement Reducer

- Sum up counts for each letter
- Write out the result to HDFS

## 4. Run the job

10

# 1: Configure Job

## • Job class

- Encapsulates information about a job
- Controls execution of the job

```
Job job = Job.getInstance(getConf(), "StartsWithCount");
```

## • A job is packaged within a jar file

- Hadoop Framework distributes the jar on your behalf
- Needs to know which jar file to distribute
- The easiest way to specify the jar that your job resides in is by calling job.setJarByClass

```
job.setJarByClass(getClass());
```

- Hadoop will locate the jar file that contains the provided class

11

# 1: Configure Job - Specify Input

```
TextInputFormat.addInputPath(job, new Path(args[0]));  
job.setInputFormatClass(TextInputFormat.class);
```

- **Can be a file, directory or a file pattern**
  - Directory is converted to a list of files as an input
- **Input is specified by implementation of InputFormat - in this case TextInputFormat**
  - Responsible for creating splits and a record reader
  - Controls input types of key-value pairs, in this case LongWritable and Text
  - File is broken into lines, mapper will receive 1 line at a time

12

## Side Node – Hadoop IO Classes

- **Hadoop uses it's own serialization mechanism for writing data in and out of network, database or files**
  - Optimized for network serialization
  - A set of basic types is provided
  - Easy to implement your own
- **org.apache.hadoop.io package**
  - LongWritable for Long
  - IntWritable for Integer
  - Text for String
  - Etc...

13



# 1: Configure Job - Specify Output

```
TextOutputFormat.setOutputPath(job, new Path(args[1]));  
job.setOutputFormatClass(TextOutputFormat.class);
```

- **OutputFormat** defines specification for outputting data from Map/Reduce job
- **Count job utilizes an implementation of OutputFormat - TextOutputFormat**
  - Define output path where reducer should place its output
    - If path already exists then the job will fail
  - Each reducer task writes to its own file
    - By default a job is configured to run with a single reducer
  - Writes key-value pair as plain text

14

# 1: Configure Job - Specify Output

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

- **Specify the output key and value types for both mapper and reducer functions**
  - Many times the same type
  - If types differ then use
    - setMapOutputKeyClass()
    - setMapOutputValueClass()

15

# 1: Configure Job

- **Specify Mapper, Reducer and Combiner**
  - At a minimum will need to implement these classes
  - Mappers and Reducer usually have same output key

```
job.setMapperClass(StartsWithCountMapper.class);  
job.setReducerClass(StartsWithCountReducer.class);  
job.setCombinerClass(StartsWithCountReducer.class);
```

16

# 1: Configure Job

- **job.waitForCompletion(true)**
  - Submits and waits for completion
  - The boolean parameter flag specifies whether output should be written to console
  - If the job completes successfully 'true' is returned, otherwise 'false' is returned

17



# Our Count Job is configured to

- **Chop up text files into lines**
- **Send records to mappers as key-value pairs**
  - Line number and the actual value
- **Mapper class is `StartsWithCountMapper`**
  - Receives key-value of `<IntWritable,Text>`
  - Outputs key-value of `<Text, IntWritable>`
- **Reducer class is `StartsWithCountReducer`**
  - Receives key-value of `<Text, IntWritable>`
  - Outputs key-values of `<Text, IntWritable>` as text
- **Combiner class is `StartsWithCountReducer`**

18

## 1: Configure Count Job

Provides Configuration support.  
More on this later...

```
public class StartsWithCountJob extends Configured implements Tool{
    @Override
    public int run(String[] args) throws Exception {
        Job job = Job.getInstance(getConf(), "StartsWithCount");
        job.setJarByClass(getClass());

        // configure output and input source
        TextInputFormat.addInputPath(job, new Path(args[0]));
        job.setInputFormatClass(TextInputFormat.class);

        // configure mapper and reducer
        job.setMapperClass(StartsWithCountMapper.class);
        job.setCombinerClass(StartsWithCountReducer.class);
        job.setReducerClass(StartsWithCountReducer.class);
```

...

...

19

# StartsWithCountJob.java Continued...

```
...  
        // configure output  
        TextOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.setOutputFormatClass(TextOutputFormat.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        return job.waitForCompletion(true) ? 0 : 1;  
    }  
  
    public static void main(String[] args) throws Exception {  
        int exitCode = ToolRunner.run(  
            new StartsWithCountJob(), args);  
        System.exit(exitCode);  
    }  
}
```

Will need an actual java main that will execute the job. More on this later...

20

## 2: Implement Mapper class

- **Class has 4 Java Generics\*\* parameters**
  - (1) input key (2) input value (3) output key (4) output value
  - Input and output utilizes hadoop's IO framework
    - org.apache.hadoop.io
- **Your job is to implement map() method**
  - Input key and value
  - Output key and value
  - Logic is up to you
- **map() method injects Context object, use to:**
  - Write output
  - Create your own counters

\*\*Java Generics provide a mechanism to abstract Java types. To learn more visit <http://docs.oracle.com/javase/tutorial/extra/generics/index.html>

21

## 2: Implement Mapper

```
public class StartsWithCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable countOne = new IntWritable(1);  
    private final Text reusableText = new Text();
```

```
    @Override
```

```
    protected void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {
```

```
        StringTokenizer tokenizer = new StringTokenizer(value.toString());  
        while (tokenizer.hasMoreTokens()) {  
            reusableText.set(tokenizer.nextToken().substring(0, 1));  
            context.write(reusableText, countOne);  
        }
```

```
    }
```

Input key and value



Output key and value



22

## 3: Implement Reducer

- **Analogous to Mapper – generic class with four types**
  - (1) input key (2) input value (3) output key (4) output value
  - The output types of map functions must match the input types of reduce function
    - In this case Text and IntWritable
  - Map/Reduce framework groups key-value pairs produced by mapper by key
    - For each key there is a set of one or more values
    - Input into a reducer is sorted by key
    - Known as Shuffle and Sort
  - Reduce function accepts key->setOfValues and outputs key-value pairs
    - Also utilizes Context object (similar to Mapper)

23

## 3: Implement Reducer

```
public class StartsWithCountReducer extends
    Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    protected void reduce(Text token,
        Iterable<IntWritable> counts,
        Context context) throws IOException, InterruptedException {
        int sum = 0;

        for (IntWritable count : counts) {
            sum += count.get();
        }
        context.write(token, new IntWritable(sum));
    }
}
```

Input key and a set of corresponding values

Produce key-value pairs

24

## 3: Reducer as a Combiner

- **Combine data per Mapper task to reduce amount of data transferred to reduce phase**
- **Reducer can very often serve as a combiner**
  - Only works if reducer's output key-value pair types are the same as mapper's output types
- **Combiners are not guaranteed to run**
  - Optimization only
  - Not for critical logic
- **More about combiners later**

25

## 4: Run Count Job

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar \
  mr.wordcount.StartsWithCountJob \
  /training/data/hamlet.txt \
  /training/playArea/wordCount/
```

Job's class  
Input file  
Output directory

```
....
....
2012-05-15 18:03:25,372 INFO mapreduce.Job (Job.java:submit(1225)) - The url to track the job: http://hadoop-
laptop:8088/proxy/application_1336894075975_0011/
2012-05-15 18:03:25,373 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1270)) - Running job:
job_1336894075975_0011
2012-05-15 18:03:31,939 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1291)) - Job job_1336894075975_0011
running in uber mode : false
2012-05-15 18:03:31,941 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 0% reduce 0%
2012-05-15 18:03:45,056 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 33% reduce 0%
2012-05-15 18:03:48,082 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 38% reduce 0%
2012-05-15 18:03:57,131 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 52% reduce 0%
2012-05-15 18:04:00,177 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 72% reduce 0%
2012-05-15 18:04:03,194 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 72% reduce 16%
2012-05-15 18:04:09,230 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 78% reduce 16%
2012-05-15 18:04:12,244 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 82% reduce 16%
2012-05-15 18:04:21,292 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 83% reduce 16%
2012-05-15 18:04:22,312 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 100% reduce 16%
2012-05-15 18:04:23,324 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 100% reduce 100%
2012-05-15 18:04:23,329 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1309)) - Job job_1336894075975_0011
completed successfully
2012-05-15 18:04:23,464 INFO mapreduce.Job (Job.java:monitorAndPrintJob(1316)) - Counters: 44
File System Counters
FILE: Number of bytes read=1010922
FILE: Number of bytes written=1494114
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=97833472
....
```

26

## Output From Your Job

- **Provides job id**

```
INFO mapreduce.Job (Job.java:monitorAndPrintJob(1270)) -
Running job: job_1337387252832_0002
```

- Used to identify, monitor and manage the job

- **Shows number of generated splits**

```
mapreduce.JobSubmitter (JobSubmitter.java:submitJobInternal(362))
- number of splits:1
```

- **Reports the Progress**

```
INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 0% reduce 0%
INFO mapreduce.Job (Job.java:monitorAndPrintJob(1298)) - map 100% reduce 0%
```

- **Displays Counters – statistics for the job**

- Sanity check that the numbers match what you expected

27


## 4: Run Count Job - Continued

...  
...

Map-Reduce Framework

```
Map input records=5158
Map output records=34189
Map output bytes=205134
Map output materialized bytes=558
Input split bytes=115
Combine input records=34189
Combine output records=69
Reduce input groups=69
Reduce shuffle bytes=558
Reduce input records=69
Reduce output records=69
Spilled Records=138
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=62
```

...  
...



28

## Output of Count Job

```
$ hdfs dfs -cat /training/playArea/wordCount/part-r-00000 | more
```

```
"      8
#      1
&      1
'     185
(     86
*     16
-      1
.      4
/      1
1     29
2      7
3      2
4      1
6      2
9      1
<     12
?      2
A     722
B     325
...
```

- **Output is written to the configured output directory**
  - /training/playArea/wordCount/
- **One output file per Reducer**
  - part-r-xxxxxx format
- **Output is driven by TextOutputFormat class**

29



# \$yarn command

- **yarn script with a class argument command launches a JVM and executes the provided Job**

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar \  
    mr.wordcount.StartsWithCountJob      \  
    /training/playArea/hamlet.txt         \  
    /training/playArea/wordCount/
```

- **You could use straight java but yarn script is more convenient**
  - Adds hadoop's libraries to CLASSPATH
  - Adds hadoop's configurations to Configuration object
    - Ex: core-site.xml, mapred-site.xml, \*.xml
  - You can also utilize \$HADOOP\_CLASSPATH environment variable

30

© 2012 [coreservlets.com](http://coreservlets.com) and [Dima May](#)



## Wrap-Up

**Customized Java EE Training:** <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.

Developed and taught by well-known author and developer. At public venues or onsite at *your* location.

# Summary

- **In this lecture we**
  - Wrote a MapReduce Job
  - Implemented Map and Reduce Functions
  - Executed the job and analyzed the output

32

© 2012 [coreservlets.com](http://coreservlets.com) and [Dima May](#)



## Questions?

[JSF 2, PrimeFaces, Java 7, Ajax, jQuery, Hadoop, RESTful Web Services, Android, Spring, Hibernate, Servlets, JSP, GWT, and other Java EE training.](#)

**Customized Java EE Training:** <http://courses.coreservlets.com/>  
Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.  
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.