



Map Reduce 2.0



Input and Output

Originals of Slides and Source Code for Examples:

<http://www.coreservlets.com/hadoop-tutorial/>

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.



**For live Hadoop training, please see courses
at <http://courses.coreservlets.com/>.**

**Taught by the author of this Hadoop tutorial. Available
at public venues, or customized versions can be held
on-site at your organization.**

- Courses developed and taught by Marty Hall
 - JSF 2, PrimeFaces, servlets/JSP, Ajax, jQuery, Android development, Java 6 or 7 programming, custom mix of topics
 - Ajax courses can concentrate on 1 library (jQuery, Prototype/Scriptaculous, Ext-JS, Dojo, etc.) or survey several
 - Courses developed and taught by coreservlets.com experts (edited by Marty)
 - **Hadoop**, Spring, Hibernate/JPA, GWT, SOAP-based and RESTful Web Services
- Contact hall@coreservlets.com for details**

Agenda

- **MapReduce Theory**
- **Types of Keys and Values**
- **Input and Output Formats**
- **Discuss Anatomy of**
 - Mappers
 - Reducers
 - Combiners
 - Partitioners

4

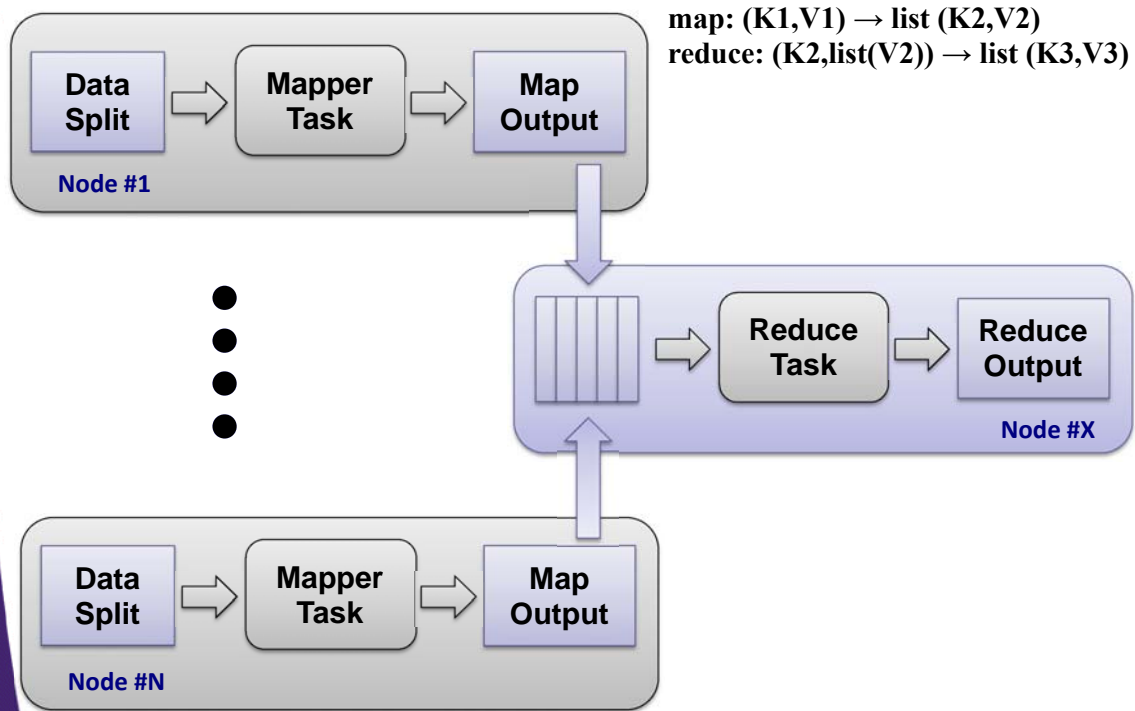
MapReduce Theory

- **Map and Reduce functions produce input and output**
 - Input and output can range from Text to Complex data structures
 - Specified via Job's configuration
 - Relatively easy to implement your own
- **Generally we can treat the flow as**

$$\text{map: } (K1, V1) \rightarrow \text{list } (K2, V2)$$
$$\text{reduce: } (K2, \text{list}(V2)) \rightarrow \text{list } (K3, V3)$$
 - Reduce input types are the same as map output types

5

Map Reduce Flow of Data



6

Key and Value Types

- **Utilizes Hadoop's serialization mechanism for writing data in and out of network, database or files**
 - Optimized for network serialization
 - A set of basic types is provided
 - Easy to implement your own
- **Extends Writable interface**
 - Framework's serialization mechanisms
 - Defines how to read and write fields
 - org.apache.hadoop.io package

7

Key and Value Types

- **Keys must implement WritableComparable interface**
 - Extends Writable and java.lang.Comparable<T>
 - Required because keys are sorted prior reduce phase
- **Hadoop is shipped with many default implementations of WritableComparable<T>**
 - Wrappers for primitives (String, Integer, etc...)
 - Or you can implement your own

8

WritableComparable<T> Implementations

Hadoop's Class	Explanation
BooleanWritable	Boolean implementation
BytesWritable	Bytes implementation
DoubleWritable	Double implementation
FloatWritable	Float implementation
IntWritable	Int implementation
LongWritable	Long implementation
NullWritable	Writable with no data

9

Implement Custom WritableComparable<T>

- **Implement 3 methods**

- write(DataOutput)
 - Serialize your attributes
- readFields(DataInput)
 - De-Serialize your attributes
- compareTo(T)
 - Identify how to order your objects
 - If your custom object is used as the key it will be sorted prior to reduce phase

10

BlogWritable – Implementation of WritableComparable<T>

```
public class BlogWritable implements
    WritableComparable<BlogWritable> {

    private String author;
    private String content;
    public BlogWritable(){}
    public BlogWritable(String author, String content) {
        this.author = author;
        this.content = content;
    }
    public String getAuthor() {
        return author;
    }

    public String getContent() {
        return content;
    }

    ...
    ...
    ...
}
```

11

BlogWritable – Implementation of WritableComparable<T>

```
...
...
@Override
public void readFields(DataInput input) throws IOException {
    author = input.readUTF();
    content = input.readUTF();
}

@Override
public void write(DataOutput output) throws IOException {
    output.writeUTF(author);
    output.writeUTF(content);
}

@Override
public int compareTo(BlogWritable other) {
    return author.compareTo(other.author);
}
}
```

1. How the data is read

2. How to write data

3. How to order BlogWritables

12

Mapper

- **Extend Mapper class**
 - Mapper<KeyIn, ValueIn, KeyOut, ValueOut>
- **Simple life-cycle**
 1. The framework first calls setup(Context)
 2. for each key/value pair in the split:
 - map(Key, Value, Context)
 3. Finally cleanup(Context) is called

13

InputSplit

- **Splits are a set of logically arranged records**
 - A set of lines in a file
 - A set of rows in a database table
- **Each instance of mapper will process a single split**
 - Map instance processes one record at a time
 - `map(k,v)` is called for each record
- **Splits are implemented by extending InputSplit class**

14

InputSplit

- **Framework provides many options for InputSplit implementations**
 - Hadoop's FileSplit
 - HBase's TableSplit
- **Don't usually need to deal with splits directly**
 - InputFormat's responsibility

15

InputFormat

- **Specification for reading data**
- **Creates Input Splits**
 - Breaks up work into chunks
- **Specifies how to read each split**
 - Divides splits into records
 - Provides an implementation of RecordReader

```
public abstract class InputFormat<K, V> {  
    public abstract  
        List<InputSplit> getSplits(JobContext context)  
                                throws IOException, InterruptedException;  
  
    public abstract  
        RecordReader<K,V> createRecordReader(InputSplit split,  
                                                TaskAttemptContext context )  
                                throws IOException, InterruptedException;  
}
```

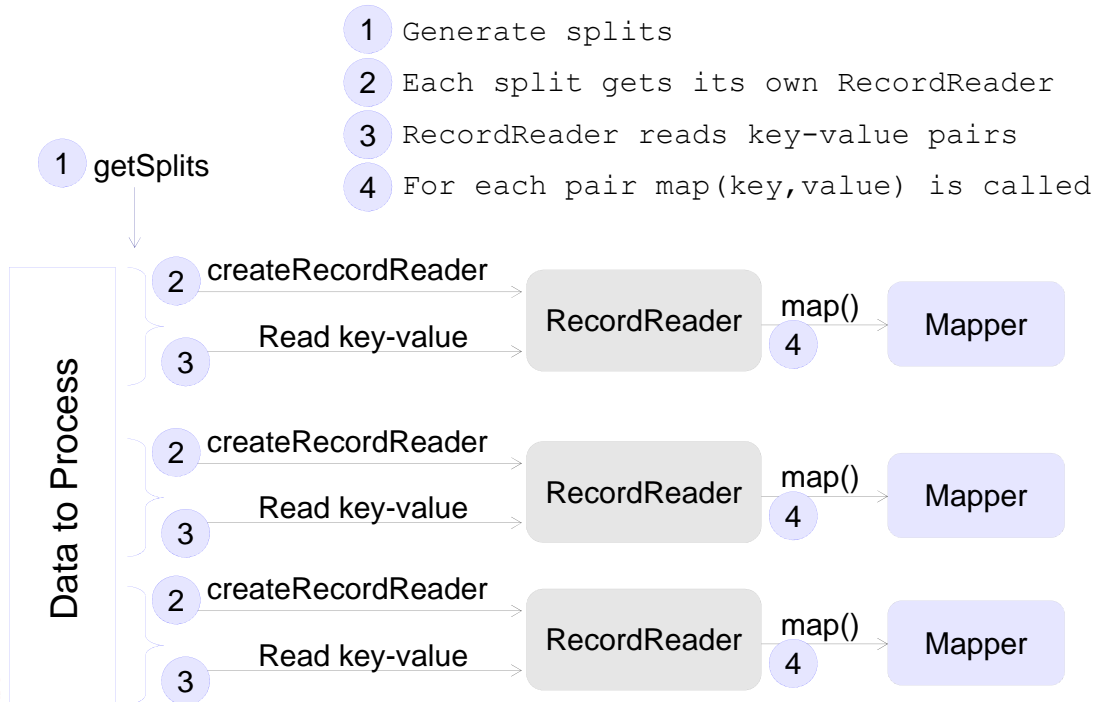
16

Framework's Usage of InputFormat Implementation

1. **Calculate splits by calling InputFormat.getSplits**
2. **For each split schedule a map task**
 - Distributed between the cluster
 - Each Map executes in its own JVM
3. **For each Mapper instance a reader is retrieved by InputFormat.createRecordReader**
 - Takes InputSplit instance as a parameter
4. **RecordReader generates key-value pairs**
5. **map() method is called for each key-value pair**

17

Framework's Usage of InputFormat Implementation



18

Hadoop's InputFormats

- **Hadoop eco-system is packaged with many InputFormats**
 - TextInputFormat
 - NLineInputFormat
 - DBInputFormat
 - TableInputFormat (HBASE)
 - StreamInputFormat
 - SequenceFileInputFormat
 - Etc...
- **Configure on a Job object**
 - `job.setInputFormatClass(XXXInputFormat.class);`

19

TextInputFormat

- **Plaint Text Input**
- **Default format**

Split:	Single HDFS block (can be configured)
Record:	Single line of text; linefeed or carriage-return used to locate end of line
Key:	LongWritable - Position in the file
Value:	Text - line of text

** Please see StartsWithCountJob for sample usage

20

NLineInputFormat

- **Same as TextInputFormat but splits equal to configured N lines**

Split:	N lines; configured via <i>mapred.line.input.format</i> or <i>NLineInputFormat.setNumLinesPerSplit(job, 100);</i>
Record:	Single line of text
Key:	LongWritable - Position in the file
Value:	Text - line of text

** Please see StartsWithCountJob_NLineInput for sample usage

21

Running TextInputFormat vs. NLineInputFormat

- Two separate runs StartsWithCountJob, one with TextInputFormat configured, next with NLineInputFormat configured
- Input is /training/playArea/hamlet.txt
 - 5159 lines
 - 206.3k

```
job.setInputFormatClass(TextInputFormat.class);
```

Job ID	Name	State	Map Progress	Maps Total
job_1338595987451_0003	StartsWithCount	RUNNING	<div><div></div></div>	1

Showing 1 to 1 of 1 entries

```
job.setInputFormatClass(NLineInputFormat.class);
```

```
NLineInputFormat.setNumLinesPerSplit(job, 100);
```

Job ID	Name	State	Map Progress	Maps Total
job_1338595987451_0002	StartsWithCount	RUNNING	<div><div></div></div>	52

Showing 1 to 1 of 1 entries

of splits

22

TableInputFormat

- Converts data in HTable to format consumable to MapReduce
- Mapper must accept proper key/values

Split:	Rows in one HBase Region (provided Scan may narrow down the result)
Record:	Row, returned columns are controlled by a provided scan
Key:	ImmutableBytesWritable
Value:	Result (HBase class)

23

StartCountJob – Input from HBase

- **Let's re-write StartWithCountJob to read input from HBase table**
 - 'HBaseSamples' table, 'count:word' family-column qualifier

```
$ hbase shell
hbase(main):005:0> scan 'HBaseSamples', {COLUMNS=>'count:word'}
ROW          COLUMN+CELL
count001     column=count:word, timestamp=1338605322765, value=Elephant
count002     column=count:word, timestamp=1338605412699, value=count
count003     column=count:word, timestamp=1338605412729, value=Updating
count004     column=count:word, timestamp=1338605412757, value=all
count005     column=count:word, timestamp=1338605412780, value=regions
count006     column=count:word, timestamp=1338605412809, value=with
count007     column=count:word, timestamp=1338605412835, value=the
count008     column=count:word, timestamp=1338605412856, value=new
count009     column=count:word, timestamp=1338605412888, value=updated
count010     column=count:word, timestamp=1338605412910, value=Done
count011     column=count:word, timestamp=1338605412933, value=seconds
count012     column=count:word, timestamp=1338605414526, value=row
12 row(s) in 0.1810 seconds
```

24

StartCountJob – Input from HBase

- 1. Re-configure Job to use HBase as input**
 - Read from table 'HBaseSamples' and column 'count:word'
 - Construct new Job 'StartWithCountJob_HBaseInput'
 - Configure Job to use new Mapper
 - New mapper now has to accept HBase Writables
 - ImmutableBytesWritable for key
 - Result for value
 - Keep reducer and combiner the same
- 2. Implement a new Mapper**
 - Grab the value from Result and write-out Text and IntWritable
 - Output is the same as in the original StartWithCountJob

25

1: Re-configure Job to use HBase as Input

```
public class StartWithCountJob_HBaseInput
    extends Configured implements Tool {

    protected final static String TABLE_NAME = "HBaseSamples";
    protected final static byte[] FAMILY = toBytes("count");
    protected final static byte[] COLUMN = toBytes("word");

    @Override
    public int run(String[] args) throws Exception {
        Job job = Job.getInstance(getConf(),
            "StartsWithCount-FromHBase");
        job.setJarByClass(getClass());
        ...
        ...
        ...
    }
}
```

Define table name, family
and column to read from

26

1: Re-configure Job to use HBase as Input

```
...
...
job.setInputFormatClass(TableInputFormat.class);
job.setMapperClass(StartsWithCountMapper_HBase.class);

Add HBase configuration and tell
framework to add HBase jars to CLASSPATH

Configuration conf = job.getConfiguration();
HBaseConfiguration.merge(conf,
    HbaseConfiguration.create(conf));
TableMapReduceUtil.addDependencyJars(job);

Specify table and
column to read from

conf.set(TableInputFormat.INPUT_TABLE, TABLE_NAME);
conf.set(TableInputFormat.SCAN_COLUMNS, "count:word");
...
...
```

27

1: Re-configure Job to use HBase as Input

...
...

Keep output the same – same reducer, combiner and output keys

```
// configure mapper and reducer
job.setCombinerClass(StartsWithCountReducer.class);
job.setReducerClass(StartsWithCountReducer.class);

// configure output
TextOutputFormat.setOutputPath(job, new Path(args[0]));
job.setOutputFormatClass(TextOutputFormat.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

return job.waitForCompletion(true) ? 0 : 1;
}

public static void main(String[] args) throws Exception {
    int exitCode = ToolRunner.run(new
        StartsWithCountJob_HBaseInput(), args);
    System.exit(exitCode);
}
}
```

28

2: Implement a New Mapper

```
public class StartsWithCountMapper_HBase
    extends TableMapper<Text, IntWritable> {
```

Extends TableMapper which assumes types for input key and value

```
private final static IntWritable countOne = new IntWritable(1);
private final Text reusableText = new Text();
```

@Override

```
protected void map(ImmutableBytesWritable key,
    Result value, Context context)
    throws IOException, InterruptedException {
```

```
    byte[] bytes = value.getValue(toBytes(FAMILY),
                                   toBytes(COLUMN));
    String str = Bytes.toString(bytes);
```

Retrieve value from input column

```
    reusableText.set(str.substring(0, 1));
    context.write(reusableText, countOne);
```

```
}
```

```
}
```

29

Run StartWithCountJob_HBaseInput

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar  
mr.wordcount.StartWithCountJob_HBaseInput /training/playArea/wordCount
```

```
$ hdfs dfs -cat /training/playArea/wordCount/part-r-00000
```

```
D      1  
E      1  
U      1  
a      1  
c      1  
n      1  
r      2  
s      1  
t      1  
u      1  
w      1
```

Reducer's output



30

Combiner

- **Runs on output of map function**
- **Produces output for reduce function**

```
map: (K1,V1) → list (K2,V2)
```

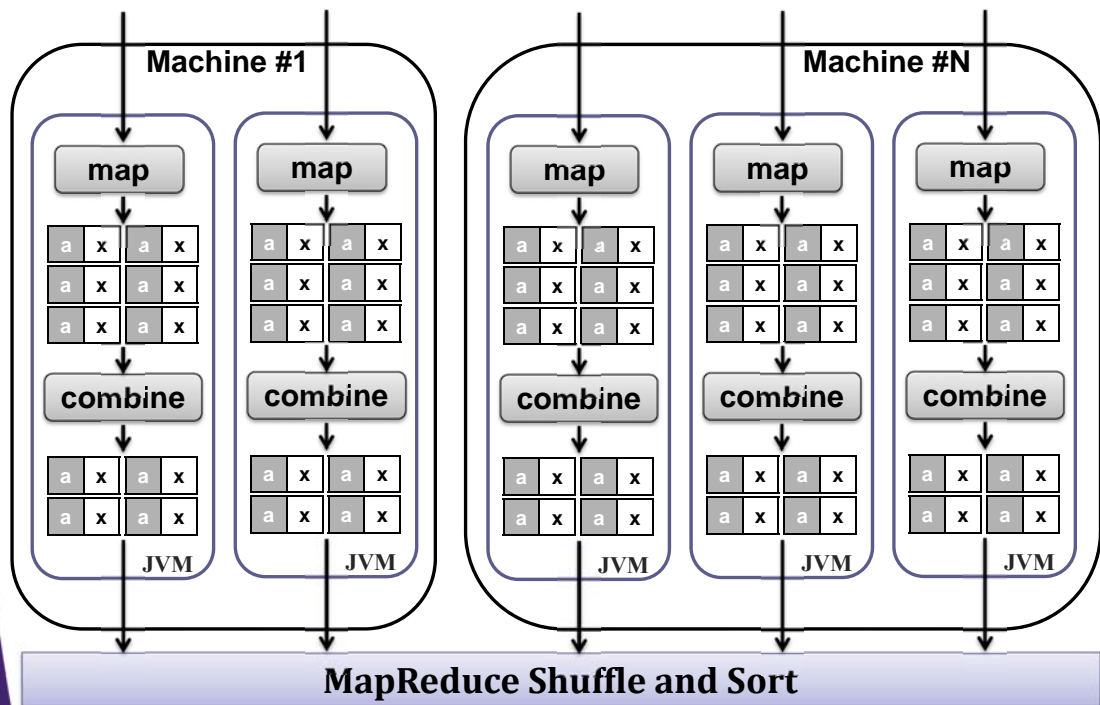
```
combine: (K2,list(V2)) → list (K2,V2)
```

```
reduce: (K2,list(V2)) → list (K3,V3)
```

- **Optimization to reduce bandwidth**
 - NO guarantees on being called
 - Maybe only applied to a sub-set of map outputs
- **Often is the same class as Reducer**
- **Each combine processes output from a single split**

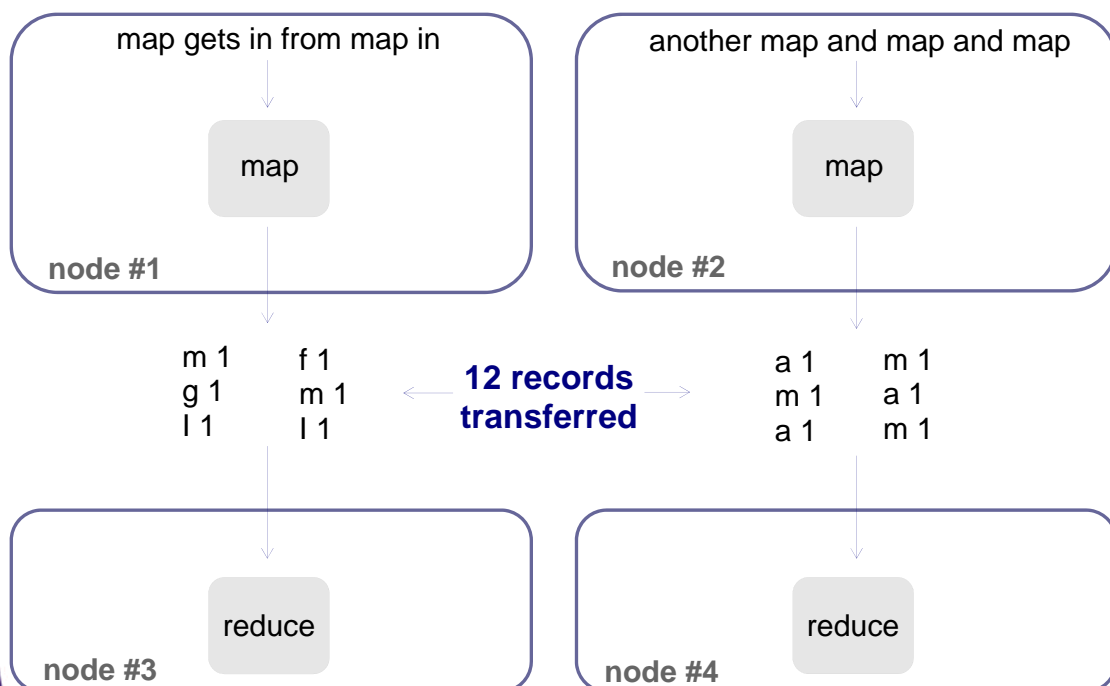
31

Combiner Data Flow




32

Sample StartsWithCountJob Run without Combiner



33

Sample StartsWithCountJob Run with Combiner



34

Specify Combiner Function

- To implement Combiner extend Reducer class
- Set combiner on Job class
 - `job.setCombinerClass(StartsWithCountReducer.class);`

35

Reducer

- **Extend Reducer class**
 - `Reducer<KeyIn, ValueIn, KeyOut, ValueOut>`
 - `KeyIn` and `ValueIn` types must match output types of mapper
- **Receives input from mappers' output**
 - Sorted on key
 - Grouped on key of key-values produced by mappers
 - Input is directed by `Partitioner` implementation
- **Simple life-cycle – similar to Mapper**
 - The framework first calls `setup(Context)`
 - for each key \rightarrow `list(value)` calls
 - `reduce(Key, Values, Context)`
 - Finally `cleanup(Context)` is called

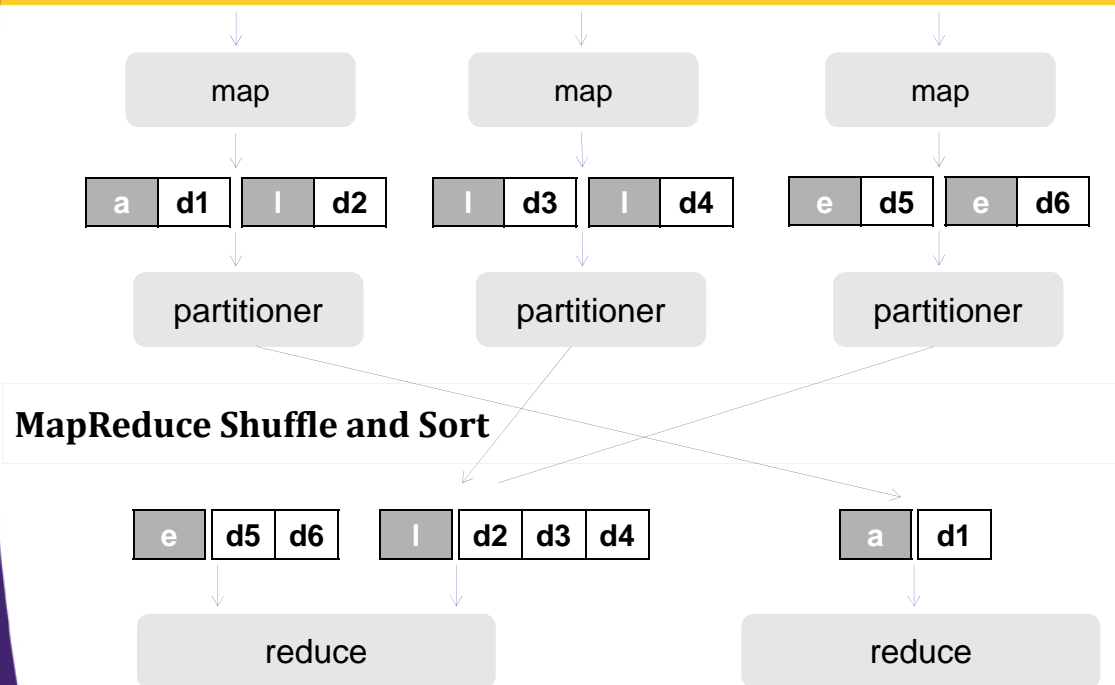
36

Reducer

- **Can configure more than 1 reducer**
 - `job.setNumReduceTasks(10);`
 - `mapreduce.job.reduces` property
 - `job.getConfiguration().setInt("mapreduce.job.reduces", 10)`
- **Partitioner implementation directs key-value pairs to the proper reducer task**
 - A partition is processed by a reduce task
 - # of partitions = # of reduce tasks
 - Default strategy is to hash key to determine partition implemented by `HashPartitioner<K, V>`

37

Partitioner Data Flow



38

Source: Jimmy Lin, Chris Dyer. [Data-Intensive Text Processing with MapReduce](#). Morgan & Claypool. 2010

HashPartitioner

```
public class HashPartitioner<K, V> extends Partitioner<K, V> {
    public int getPartition(K key, V value, int numReduceTasks) {
        return (key.hashCode() & Integer.MAX_VALUE) % numReduceTasks;
    }
}
```

- **Calculate Index of Partition:**
 - Convert key's hash into non-negative number
 - Logical AND with maximum integer value
 - Modulo by number of reduce tasks
- **In case of more than 1 reducer**
 - Records distributed evenly across available reduce tasks
 - Assuming a good hashCode() function
 - Records with same key will make it into the same reduce task
 - Code is independent from the # of partitions/reducers specified

39

Custom Partitioner

```
public class CustomPartitioner
    extends Partitioner<Text, BlogWritable>{
    @Override
    public int getPartition(Text key, BlogWritable blog,
        int numReduceTasks) {
        int positiveHash =
            blog.getAuthor().hashCode() & Integer.MAX_VALUE;

        return positiveHash % numReduceTasks;
    }
}
```

Use author's hash only, AND with
max integer to get a positive value

Assign a reduce task by index

- **All blogs with the same author will end up in the same reduce task**

40

OutputFormat

- **Specification for writing data**
 - The other side of InputFormat
- **Implementation of OutputFormat<K,V>**
- **TextOutputFormat is the default implementation**
 - Output records as lines of text
 - Key and values are tab separated “Key /t value”
 - Can be configured via
“mapreduce.output.textoutputformat.separator” property
 - Key and Value may of any type - call .toString()

41

OutputFormat

- **Validates output specification for that job**
 - You may have seen annoying messages that output directory already exists
- **Creates implementation of RecordWriter**
 - Responsible for actually writing data
- **Creates implementation of OutputCommitter**
 - Set-up and clean-up Job's and Tasks' artifacts (ex. Directories)
 - Commit or discard tasks output

42

OutputFormat Interface

```
public abstract class OutputFormat<K, V> {  
  
    public abstract RecordWriter<K, V>  
        getRecordWriter(TaskAttemptContext context  
                        ) throws IOException, InterruptedException;
```

Provide implementation of
RecordReader

```
    public abstract void checkOutputSpecs(JobContext context  
                                        ) throws IOException, InterruptedException;
```

Provide
implementation of
OutputCommitter

Validates output specification for
that job

```
    public abstract  
        OutputCommitter getOutputCommitter(TaskAttemptContext context  
                                        ) throws IOException, InterruptedException;
```

```
}
```

43

Hadoop's OutputFormats

- **Hadoop eco-system is packaged with many OutputFormats**
 - TextOutputFormat
 - DBOutputFormat
 - TableOutputFormat (HBASE)
 - MapFileOutputFormat
 - SequenceFileOutputFormat
 - NullOutputFormat
 - Etc...
- **Configure on Job object**
 - `job.setOutputFormatClass(XXXOutputFormat.class);`
 - `job.setOutputKeyClass(XXXKey.class);`
 - `job.setOutputValueClass(XXXValue.class);`

44

TextOutputFormat

- **Outputs plain text**
- **Saves key-value pairs separated by tab**
 - Configured via
`mapreduce.output.textoutputformat.separator` property
- **Set output path**
 - `TextOutputFormat.setOutputPath(job, new Path(myPath));`

45

TableOutputFormat

- Saves data into HTable
- Reducer output key is ignored
- Reducer output value must be HBase's Put or Delete objects

46

TableOutputFormat

```
public class StartWithCountJob_HBase
    extends Configured implements Tool {

    protected final static String TABLE_NAME = "HBaseSamples";
    protected final static String FAMILY = "count";
    protected final static String INPUT_COLUMN = "word";

    protected final static String RESULT_COLUMN = "result";

    @Override
    public int run(String[] args) throws Exception {
        Job job =
            Job.getInstance(getConf(), "StartsWithCount-HBase");
        job.setJarByClass(getClass());
        ...
        ...
    }
}
```

Input and output column resides
Under the same family (doesn't have to)

Name the job

47

TableOutputFormat

```
Scan scan = new Scan();
scan.addColumn(toBytes(FAMILY), toBytes(INPUT_COLUMN));
TableMapReduceUtil.initTableMapperJob(
    TABLE_NAME,
    scan,
    StartsWithCountMapper_HBase.class,
    Text.class,
    IntWritable.class,
    job);

TableMapReduceUtil.initTableReducerJob(
    TABLE_NAME,
    StartsWithCountReducer_HBase.class,
    job);
job.setNumReduceTasks(1);
return job.waitForCompletion(true) ? 0 : 1;
}

public static void main(String[] a) throws Exception {
    int code = ToolRunner.run(new StartWithCountJob_HBase(), a);
    System.exit(code);
}
}
```

Utilize HBase's TableMapReduceUtil to setup the job; internally delegate to job.XXX methods

48

TableOutputFormat

```
public class StartsWithCountReducer_HBase extends
    TableReducer<Text, IntWritable, ImmutableBytesWritable> {

    @Override
    protected void reduce(Text key, Iterable<IntWritable> counts,
        Context context)
        throws IOException, InterruptedException {
        int sum = 0;

        for (IntWritable count : counts) {
            sum += count.get();
        }

        Put put = new Put(key.copyBytes());
        put.add(toBytes(FAMILY), toBytes(RESULT_COLUMN),
            toBytes(Integer.toString(sum)));
        context.write(null, put);
    }
}
```

Reducer must output either Put or Delete object

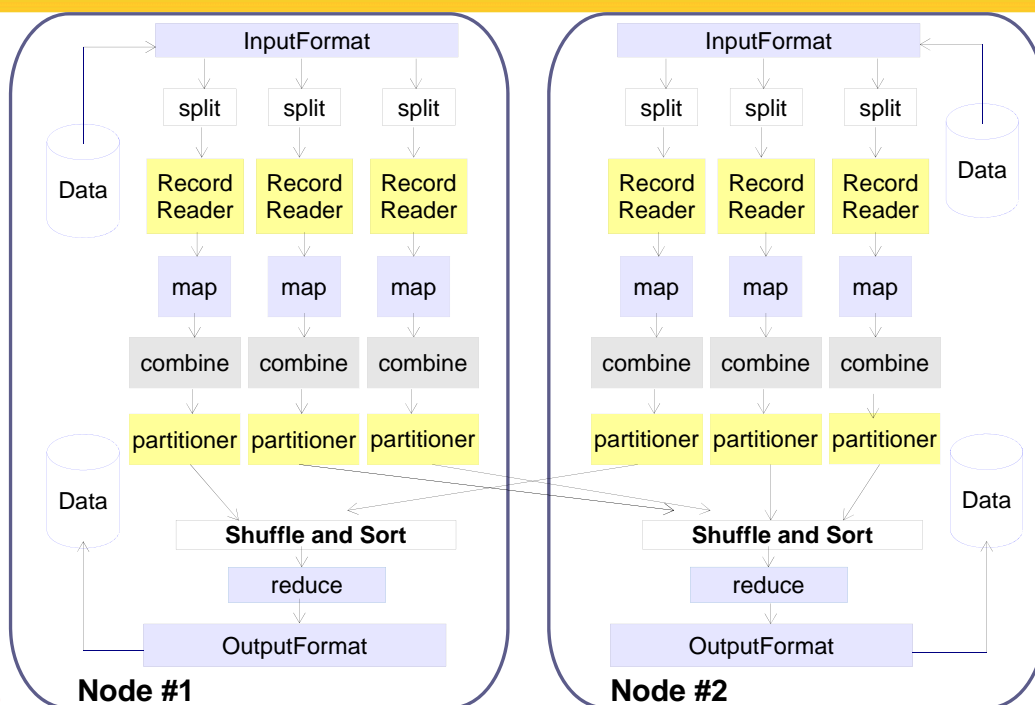
49

TableOutputFormat

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar mr.wordcount.StartWithCountJob_HBase
$ hbase shell
hbase(main):029:0> scan 'HBaseSamples', {COLUMN=>'count:result'}
ROW          COLUMN+CELL
D            column=count:result, timestamp=1338951024384, value=1
E            column=count:result, timestamp=1338951024384, value=1
U            column=count:result, timestamp=1338951024384, value=1
a            column=count:result, timestamp=1338951024384, value=1
c            column=count:result, timestamp=1338951024386, value=1
n            column=count:result, timestamp=1338951024386, value=1
r            column=count:result, timestamp=1338951024386, value=2
s            column=count:result, timestamp=1338951024386, value=1
t            column=count:result, timestamp=1338951024386, value=1
u            column=count:result, timestamp=1338951024386, value=1
w            column=count:result, timestamp=1338951024386, value=1
12 row(s) in 0.0530 seconds
```

50

Component Overview



51

Source: Yahoo! Inc. "Hadoop Tutorial from Yahoo!". 2012



Wrap-Up

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.

Summary

- **In this lecture we learned about**
 - MapReduce Theory
 - Types of Keys and Values
 - Input and Output Formats
 - Anatomy of
 - Mappers
 - Reducers
 - Combiners
 - Partitioners



Questions?

[JSF 2, PrimeFaces, Java 7, Ajax, jQuery, Hadoop, RESTful Web Services, Android, Spring, Hibernate, Servlets, JSP, GWT, and other Java EE training.](#)

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.