




Map Reduce Features

Originals of Slides and Source Code for Examples:

<http://www.coreservlets.com/hadoop-tutorial/>

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.



**For live Hadoop training, please see courses
at <http://courses.coreservlets.com/>.**

**Taught by the author of this Hadoop tutorial. Available
at public venues, or customized versions can be held
on-site at your organization.**

- Courses developed and taught by Marty Hall
 - JSF 2, PrimeFaces, servlets/JSP, Ajax, jQuery, Android development, Java 6 or 7 programming, custom mix of topics
 - Ajax courses can concentrate on 1 library (jQuery, Prototype/Scriptaculous, Ext-JS, Dojo, etc.) or survey several
 - Courses developed and taught by coreservlets.com experts (edited by Marty)
 - **Hadoop**, Spring, Hibernate/JPA, GWT, SOAP-based and RESTful Web Services
- Contact hall@coreservlets.com for details**

Agenda

- **Counters**
- **Speculative Execution**
- **Distributed Cache**

4

Counters

- **Instrument Job's metrics**
 - Gather statistics
 - Quality control – confirm what was expected
 - Diagnostics
- **Framework provides a set of built-in metrics**
 - For example bytes processed for input and output
- **User can create new counters**
 - Number of records consumed
 - Number of errors or warnings
- **Counters are divided into groups**
- **Tracks total, mapper and reducer counts**

5

Built-in Counters

- Maintains and sums up counts
- Several groups for built-in counters
 - Job Counters – documents number of map and reduce tasks launched, number of failed tasks
 - File System Counters – number of bytes read and written
 - Map-Reduce Framework – mapper, reducer, combiner input and output records counts, time and memory statistics

6

Job Counters

- Web UI exposes counters for each Job



Counters for job_1339260384899_0001

| Counter Group | Name | Map | Reduce | Total |
|----------------------|---|----------|--------|----------|
| File System Counters | FILE: Number of bytes read | 298 | 0 | 298 |
| | FILE: Number of bytes written | 104542 | 0 | 104542 |
| | FILE: Number of large read operations | 0 | 0 | 0 |
| | FILE: Number of read operations | 0 | 0 | 0 |
| | FILE: Number of write operations | 0 | 0 | 0 |
| | HDFS: Number of bytes read | 51853 | 0 | 51853 |
| | HDFS: Number of bytes written | 62586 | 0 | 62586 |
| | HDFS: Number of large read operations | 0 | 0 | 0 |
| | HDFS: Number of read operations | 37 | 0 | 37 |
| | HDFS: Number of write operations | 14 | 0 | 14 |
| Job Counters | Data-local map tasks | 0 | 0 | 1 |
| | Launched map tasks | 0 | 0 | 1 |
| | Total time spent by all maps in occupied slots (ms) | 0 | 0 | 4408 |
| | | | | |
| Map-Reduce Framework | CPU time spent (ms) | 1920 | 0 | 1920 |
| | Failed Shuffles | 0 | 0 | 0 |
| | GC time elapsed (ms) | 21 | 0 | 21 |
| | Input split bytes | 169 | 0 | 169 |
| | Map input records | 1 | 0 | 1 |
| | Map output records | 0 | 0 | 0 |
| | Merged Map outputs | 0 | 0 | 0 |
| | Physical memory (bytes) snapshot | 85540864 | 0 | 85540864 |
| | Spilled Records | 0 | 0 | 0 |
| | Total committed heap usage (bytes) | 37355520 | 0 | 37355520 |
| | | | | |

7

User Defined Counters

- **You can create new counters**
- **Increment counters in Reducer and/or Mapper classes**
 - Framework accurately sums up counts between various stages and produces totals
 - Accounts for failed tasks

8

Implement User-Defined Counters

- 1. Retrieve Counter from Context object**
 - Framework injects Context object into map and reduce methods
- 2. Increment Counter's value**
 - Can increment by 1 or more

9

1: Retrieve Counter from Context

- **Utilize Context object**
 - void map(Key key, Value value, Context context)
 - void reduce(Key key, Iterable<Value> values, Context context)
- **Although map's and reduce's Context type is not the same they both extend from**

`org.apache.hadoop.mapreduce.TaskAttemptContext`

- **TaskAttemptContext provides two ways to retrieve counters**
 - public Counter getCounter(String groupName, String counterName);
 - public Counter getCounter(Enum<?> counterName);
 - Figures out group name from fully qualified classname of enum - `enum.getDeclaringClass().getName()`

10

2: Increment Counter's Value

- **Increment or even set the value**
 - void setValue(long value);
 - void increment(long incr);

11

StartsWithCountMapper with Counters

- Recall the StartsWithCountJob
- Update Mapper to document counts for
 - Total tokens processed
 - Number of tokens that start with uppercase
 - Number of tokens that start with lowercase
- First create an enum to reference these counters:

```
public enum Tokens {  
    Total, FirstCharUpper, FirstCharLower  
}
```

12

StartsWithCountMapper_UserCounters.java

```
@Override  
protected void map(LongWritable key, Text value, Context context)  
    throws IOException, InterruptedException {  
  
    StringTokenizer tokenizer = new StringTokenizer(value.toString());  
    while (tokenizer.hasMoreTokens()) {  
        String token = tokenizer.nextToken();  
        reusableText.set(token.substring(0, 1));  
        context.write(reusableText, countOne);  
  
        context.getCounter(Tokens.Total).increment(1);  
        char firstChar = token.charAt(0);  
  
        if (Character.isUpperCase(firstChar)) {  
            context.getCounter(Tokens.FirstCharUpper).increment(1);  
        } else {  
            context.getCounter(Tokens.FirstCharLower).increment(1);  
        }  
    }  
}
```

Keep count of total tokens processed

Stats on tokens that start with upper case vs. lowercase

13

Run StartsWithCountMapper_UserCounters

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar  
mr.wordcount.StartsWithCountJob_UserCounters  
/training/data/hamlet.txt /training/playArea/wordCount/  
...  
...  
...  
Map output records=34189  
Map output bytes=205134  
Combine input records=34189  
Combine output records=69  
Reduce input records=69  
Reduce output records=69  
mr.wordcount.StartsWithCountMapper_UserCounters$Tokens  
FirstCharLower=26080  
FirstCharUpper=8109  
Total=34189  
File Input Format Counters  
Bytes Read=211294  
File Output Format Counters  
Bytes Written=385
```

Job configures new mapper with counts

Group was derived from the class name

Total # of tokens should match Map output records

14

Customize Counter's Names

- Can customize counter and group names when using enums
 1. Create a properties file <classname>.properties defining counter name properties
 - Inner classes are substituted by underscore
 - For example: org.com.MyMapper\$CustomEnum would be MyMapper_CustomEnum.properties
 2. Place properties file in the same package as the class that defines Enum

15

1: Create Properties File

- In our case the enum was defined in
 - mr.wordcount.StartsWithCountMapper_UserCounters\$Tokens
- Therefore the file is to be named
 - StartsWithCountMapper_UserCounters_Tokens.properties
- Define Group and Counter names:

```
CounterGroupName = Token Processed
Total.name=Total Tokens Processed
FirstCharUpper.name=Tokens start with Uppercase
FirstCharLower.name=Tokens start with Lowercase
```

16

Test Counter Re-Naming

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar
mr.wordcount.StartsWithCountJob_UserCounters
/training/data/hamlet.txt /training/playArea/wordCount/
...
...
...
    Map output records=34189
    Map output bytes=205134
    Combine input records=34189
    Combine output records=69
    Reduce input records=69
    Reduce output records=69
Token Processed
    Tokens start with Lowercase=26080
    Tokens start with Uppercase=8109
    Total Tokens Processed =34189
File Input Format Counters
    Bytes Read=211294
File Output Format Counters
    Bytes Written=385
```

New Names are mapped by the framework

17

Retrieving Counters

1. Web-UI

| Counter Group | Name | Counters | Map | Reduce | Total |
|----------------------|--|----------|-----|--------|-------|
| Job | File, Number of bytes read | 298 | 0 | 298 | |
| | File, Number of bytes written | 204542 | 0 | 204542 | |
| | File, Number of bytes read operations | 0 | 0 | 0 | |
| | File, Number of bytes written operations | 0 | 0 | 0 | |
| | File, Number of bytes read operations | 0 | 0 | 0 | |
| File System Counters | File, Number of bytes read | 10883 | 0 | 10883 | |
| | File, Number of bytes written | 42586 | 0 | 42586 | |
| | File, Number of bytes read operations | 0 | 0 | 0 | |
| | File, Number of bytes written operations | 17 | 0 | 17 | |
| | File, Number of bytes read operations | 14 | 0 | 14 | |
| Job Counters | Map, Map input records | 0 | 0 | 0 | |
| | Map, Map output records | 0 | 0 | 0 | |
| | Map, Map output bytes | 0 | 0 | 0 | |
| Map Reduce Framework | Map, Map input records | 2628 | 0 | 2628 | |
| | Map, Map output records | 0 | 0 | 0 | |
| | Map, Map output bytes | 0 | 0 | 0 | |
| | Map, Map output bytes | 149 | 0 | 149 | |
| | Map, Map output bytes | 0 | 0 | 0 | |

2. Command line

- \$ mapred job -status <job_id>

3. Java API

- Further analyze counts
- Store in a database

18

Java API to Retrieve Counters

• Print all the counters after the job is done

- Snippet from StartsWithCountJob_PrintCounters.java

```
int resultCode = job.waitForCompletion(true) ? 0 : 1;

System.out.println("Job is complete! Printing Counters:");
Counters counters = job.getCounters();

for (String groupName : counters.getGroupNames()) {
    CounterGroup group = counters.getGroup(groupName);
    System.out.println(group.getDisplayName());

    for (Counter counter : group.getUnderlyingGroup()) {
        System.out.println("  " + counter.getDisplayName() +
            " = " + counter.getValue());
    }
}
```

19

Speculative Execution

- **Job is decomposed into small tasks**
- **Job is as fast as the slowest task**
- **Given 100s or even 1000s of tasks**
 - Few tasks may run very slowly (hardware issues, other activity on that machine, configuration, etc...)
- **MapReduce framework strives to resolve slow running tasks by spawning the same task on a different machine**
 - Doesn't start speculative tasks immediately

20

Speculative Execution

- **Will spawn a speculative task when**
 - All the tasks have been started
 - Task has been running for an extended period of time
 - over a minute
 - Did not make significant progress as compared to the rest of the running tasks
- **After task's completion duplicates are killed**
- **Just an optimization**

21

Speculative Execution

- **Can be turned off by setting these properties to false**
 - `mapred.map.tasks.speculative.execution`
 - Turn on/off speculative execution for map phase
 - `mapred.reduce.tasks.speculative.execution`
 - Turn on/off speculative execution for reduce phase
- **When should I disable Speculative Execution?**
 - Task is outputting directly to a shared resource; then starting a duplicate task may cause unexpected results
 - Minimize cluster and bandwidth usage; duplicate tasks use up resources

22

Distributed Cache

- **A mechanism to distribute files**
- **Make them available to MapReduce task code**
- **yarn command provides several options to add distributed files**
- **Can also use Java API directly**
- **Supports**
 - Simple text files
 - Jars
 - Archives: zip, tar, tgz/tar.gz

23

Distributed Cache via \$ yarn Command

- **Update StartWithCount job to exclude specified start letters**
 1. Load a file which contains start letters to exclude onto distributed cache
 - utilize -files parameter with the yarn command
 2. Update Map code to utilize the exclude file

24

1: Load Exclude File Onto DistributedCache

```
$ cat $PLAY_AREA/data/startWithExcludeFile.txt
```

```
b  
c  
d  
e  
f  
G
```

Exclude tokens that start with these letters

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar \  
mr.wordcount.StartsWithCountJob_DistCache \  
-files $PLAY_AREA/data/startWithExcludeFile.txt \  
/training/data/hamlet.txt \  
/training/playArea/wordCount/
```

Using -files option yarn command will place the file onto DistributedCache.

25

2: Utilize Exclude File in the Map Code

```
public class StartsWithCountMapper_DistCache extends
    Mapper<LongWritable, Text, Text, IntWritable> {
    private Logger log =
        Logger.getLogger(StartsWithCountMapper_DistCache.class);
    private final static IntWritable countOne =
        new IntWritable(1);
    private final Text reusableText = new Text();

    public final static String EXCLUDE_FILE =
        "startWithExcludeFile.txt";
    private final Set<String> excludeSet = new HashSet<String>();
    ...
    ...
    ...
```

Will be able to directly
reference the file without
absolute path;
Constructs an exclude set

26

2: Utilize Exclude File in the Map Code

```
...
...
@Override
protected void setup(Context context) throws IOException,
    InterruptedException {
    FileReader reader = new FileReader(new File(EXCLUDE_FILE));
    try {
        BufferedReader bufferedReader = new BufferedReader(reader);
        String line;
        while ((line = bufferedReader.readLine()) != null) {
            excludeSet.add(line);
            log.info("Ignoring words that start with [" + line + "]);
        }
    } finally {
        reader.close();
    }
}
...
...
```

Framework takes care of all the magic;
the file is now stored locally and can
be referenced with just a name

27

2: Utilize Exclude File in the Map Code

```
...
...
@Override
protected void map(LongWritable key, Text value,
    Context context) throws IOException, InterruptedException {

    StringTokenizer tokenizer =
        new StringTokenizer(value.toString());

    while (tokenizer.hasMoreTokens()) {
        String firstLetter =
            tokenizer.nextToken().substring(0, 1);

        if (!excludeSet.contains(firstLetter)) {
            reusableText.set(firstLetter);
            context.write(reusableText, countOne);
        }
    }
}
```

28

Output of Distributed Cache Job

```
$ hdfs dfs -cat /training/playArea/wordCount/part-r-00000
...
...
[      122
-      1
a     2370
h     1724
i     1468
j      65
k     222
...
...
```

← Exclude letters are missing

29

Distributed Cache Inner-Workings

- **Accepts two types: files and archives**
 - Archives are unarchived on the local node
- **Items specified to the \$yarn command via -files, -libjars and -archives are copied to HDFS**
- **Prior to task execution these files are copied locally from HDFS**
 - Files now reside on a local disk – local cache
- **Files provided to the -libjars are appended to task's CLASSPATH**
- **Locally cached files become qualified to be deleted after all tasks utilizing cache complete**

30

Distributed Cache Inner-Workings

- **Files in the local cache are deleted after a 10GB threshold is reached**
 - Allow space for new files
 - Configured via `yarn.nodemanager.localizer.cache.target-size-mb` property
- **Local cache is stored under**
`${yarn.nodemanager.local-dirs}/usercache/$user/filecache`
 - Task code is not aware of the location
 - Symbolic link is created for each file, that's why we were able to reference a file without the path
 - `FileReader reader = new FileReader(new File(EXCLUDE_FILE))`

31

Java API - DistributedCache

- **Typically don't need to use directly**
 - Delegate to \$yarn command (with -files, -libjars or -archive options)
- **However when programmatic involvement is necessary use DistributedCache class**
 - File(s) to be cached must exist on HDFS
 - With \$yarn command framework takes care of this step for you
 - In the Tool – Place data into cache via methods on Job
 - In the Task – Retrieve the file from the cache

32

1: File(s) to be Cached Must Exist on HDFS

- **Add file to HDFS, either via command line or via FileSystem Java API**
- **Major difference from using \$yarn command where framework adds it to HDFS on your behalf**

```
$ hdfs dfs -put startWithExcludeFile.txt /training/data/
$ hdfs dfs -ls /training/data/
Found 3 items
  0 2011-12-24 11:21 /training/data/glob
 22 2011-12-20 22:12 /training/data/readMe.txt
 12 2012-06-17 16:08 /training/data/startWithExcludeFile.txt
```

33

2: In the Tool – Place Data Onto Cache via Methods on Job

StartsWithCountJob_DistCacheAPI.java

```
....
public int run(String[] args) throws Exception {
    Job job = Job.getInstance(getConf(), getClass().getSimpleName());
    job.setJarByClass(getClass());
    ...
    ...
    ...
    Path toCache = new Path("/training/data/startWithExcludeFile.txt");
    job.addCacheFile(toCache.toUri());
    job.createSymlink();

    return job.waitForCompletion(true) ? 0 : 1;
}
...
```

Add file to DistributedCache

Create symbolic links for all files in DistributedCache;
without the links you would have to use fully qualified
path, in this case "/training/data/startWithExcludeFile.txt"

34

3: In the Task – Retrieve the File From the Cache

- Same as before:

StartsWithCountMapper_DistCache.java

```
...
...
@Override
protected void setup(Context context) throws IOException,
    InterruptedException {
    FileReader reader = new FileReader(new File(EXCLUDE_FILE));
    try {
        BufferedReader bufferedReader = new BufferedReader(reader);
        String line;
        while ((line = bufferedReader.readLine()) != null) {
            excludeSet.add(line);
            log.info("Ignoring words that start with [" + line + "]);
        }
    } finally {
        reader.close();
    }
}
...
```

35



Wrap-Up

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.

Summary

- **We learned about**
 - Counters
 - Speculative Execution
 - Distributed Cache



Questions?

[JSF 2, PrimeFaces, Java 7, Ajax, jQuery, Hadoop, RESTful Web Services, Android, Spring, Hibernate, Servlets, JSP, GWT, and other Java EE training.](#)

Customized Java EE Training: <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.