



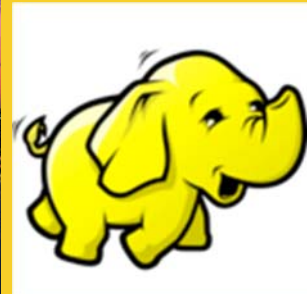
# Hive

Originals of Slides and Source Code for Examples:

<http://www.coreservlets.com/hadoop-tutorial/>

**Customized Java EE Training:** <http://courses.coreservlets.com/>

Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.  
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.



**For live Hadoop training, please see courses  
at <http://courses.coreservlets.com/>.**

**Taught by the author of this Hadoop tutorial. Available  
at public venues, or customized versions can be held  
on-site at your organization.**

- Courses developed and taught by Marty Hall
    - JSF 2, PrimeFaces, servlets/JSP, Ajax, jQuery, Android development, Java 6 or 7 programming, custom mix of topics
    - Ajax courses can concentrate on 1 library (jQuery, Prototype/Scriptaculous, Ext-JS, Dojo, etc.) or survey several
  - Courses developed and taught by [coreservlets.com](http://coreservlets.com) experts (edited by Marty)
    - **Hadoop**, Spring, Hibernate/JPA, GWT, SOAP-based and RESTful Web Services
- Contact [hall@coreservlets.com](mailto:hall@coreservlets.com) for details**

# Agenda

- **Hive Overview and Concepts**
- **Installation**
- **Table Creation and Deletion**
- **Loading Data into Hive**
- **Partitioning**
- **Bucketing**
- **Joins**



4

# Hive

- **Data Warehousing Solution built on top of Hadoop**
- **Provides SQL-like query language named HiveQL**
  - Minimal learning curve for people with SQL expertise
  - Data analysts are target audience
- **Early Hive development work started at Facebook in 2007**
- **Today Hive is an Apache project under Hadoop**
  - <http://hive.apache.org>

5

## Hive Provides

- **Ability to bring structure to various data formats**
- **Simple interface for ad hoc querying, analyzing and summarizing large amounts of data**
- **Access to files on various data stores such as HDFS and HBase**

6


## Hive

- **Hive does NOT provide low latency or real-time queries**
- **Even querying small amounts of data may take minutes**
- **Designed for scalability and ease-of-use rather than low latency responses**

7

# Hive

- **Translates HiveQL statements into a set of MapReduce Jobs which are then executed on a Hadoop Cluster**



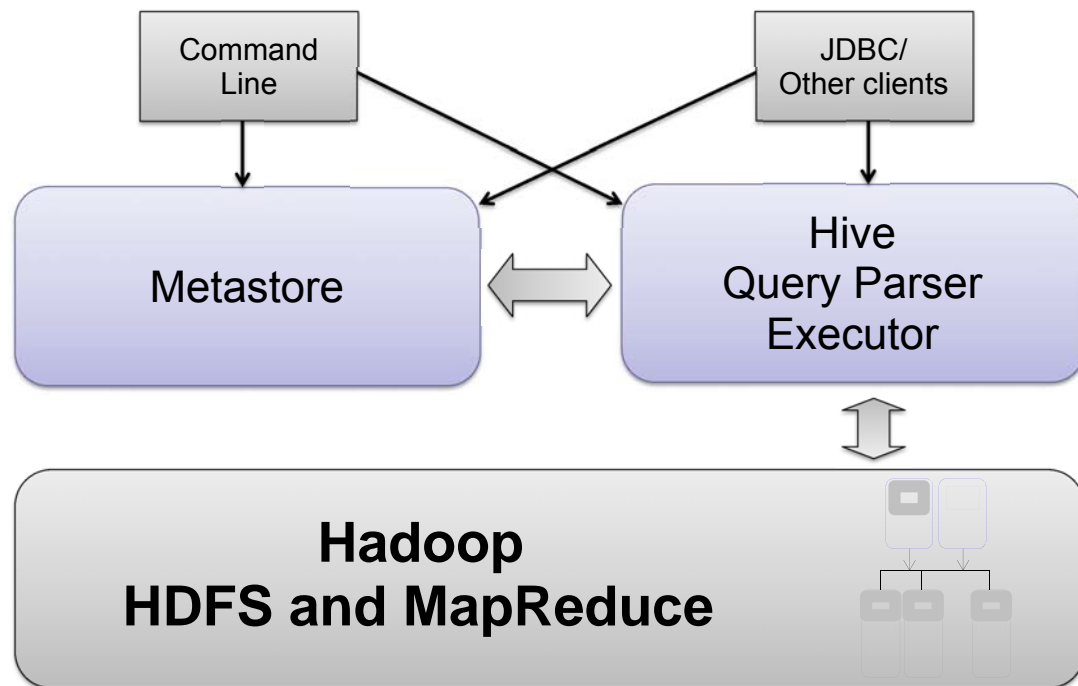
8

# Hive Metastore

- **To support features like schema(s) and data partitioning Hive keeps its metadata in a Relational Database**
  - Packaged with Derby, a lightweight embedded SQL DB
    - Default Derby based is good for evaluation and testing
    - Schema is not shared between users as each user has their own instance of embedded Derby
    - Stored in metastore\_db directory which resides in the directory that hive was started from
  - Can easily switch another SQL installation such as MySQL

9

# Hive Architecture



10

# Hive Interface Options


- **Command Line Interface (CLI)**
  - Will use exclusively in these slides
- **Hive Web Interface**
  - <https://cwiki.apache.org/confluence/display/Hive/HiveWebInterface>
- **Java Database Connectivity (JDBC)**
  - <https://cwiki.apache.org/confluence/display/Hive/HiveClient>

11

# Hive Concepts

- **Re-used from Relational Databases**

- **Database:** Set of Tables, used for name conflicts resolution
- **Table:** Set of Rows that have the same schema (same columns)
- **Row:** A single record; a set of columns
- **Column:** provides value and type for a single value



12

# Installation Prerequisites

- **Java 6**
  - Just Like Hadoop
- **Hadoop 0.20.x+**
  - No surprise here

13

# Hive Installation

- **Set \$HADOOP\_HOME environment variable**
  - Was done as a part of HDFS installation
- **Set \$HIVE\_HOME and add hive to the PATH**

```
export HIVE_HOME=$CDH_HOME/hive-0.8.1-cdh4.0.0
export PATH=$PATH:$HIVE_HOME/bin
```

- **Hive will store its tables on HDFS and those locations needs to be bootstrapped**

```
$ hdfs dfs -mkdir /tmp
$ hdfs dfs -mkdir /user/hive/warehouse
$ hdfs dfs -chmod g+w /tmp
$ hdfs dfs -chmod g+w /user/hive/warehouse
```

14

# Hive Installation

- **Similar to other Hadoop's projects Hive's configuration is in \$HIVE\_HOME/conf/hive-site.xml**

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>

  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:10040</value>
  </property>

</configuration>
```

Specify the location of ResourceManager so Hive knows where to execute MapReduce Jobs; by default Hive utilizes LocalJobRunner

15



# Run Hive

- **HDFS and YARN need to be up and running**

```
$ hive  
Hive history file=/tmp/hadoop/hive_job_log_hadoop_201207312052_1402761030.txt  
hive>
```

Hive's Interactive Command Line Interface (CLI)



16

# Simple Example

- 1. Create a Table**
- 2. Load Data into a Table**
- 3. Query Data**
- 4. Drop the Table**

17



# 1: Create a Table

- Let's create a table to store data from `$PLAY_AREA/data/user-posts.txt`

Launch Hive Command Line Interface (CLI)

```
$ cd $PLAY_AREA
```

Location of the session's log file

```
$ hive
```

Hive history file=/tmp/hadoop/hive\_job\_log\_hadoop\_201208022144\_2014345460.txt

```
hive> !cat data/user-posts.txt;
```

Can execute local commands within CLI, place a command in between ! and ;

```
user1,Funny Story,1343182026191
user2,Cool Deal,1343182133839
user4,Interesting Post,1343182154633
user5,Yet Another Blog,13431839394
hive>
```

Values are separate by ',' and each row represents a record; first value is user name, second is post content and third is timestamp

18

# 1: Create a Table

```
hive> CREATE TABLE posts (user STRING, post STRING, time BIGINT)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE;
```

1<sup>st</sup> line: creates a table with 3 columns  
2<sup>nd</sup> and 3<sup>rd</sup> line: how the underlying file should be parsed  
4<sup>th</sup> line: how to store data

Statements must end with a semicolon and can span multiple rows

```
OK
Time taken: 10.606 seconds
```

```
hive> show tables;
```

Display all of the tables

```
OK
posts
Time taken: 0.221 seconds
```

Result is displayed between "OK" and "Time taken..."

```
hive> describe posts;
```

Display schema for posts table

```
OK
user      string
post      string
time      bigint
Time taken: 0.212 seconds
```

19

## 2: Load Data Into a Table

```
hive> LOAD DATA LOCAL INPATH 'data/user-posts.txt'
> OVERWRITE INTO TABLE posts;
```

```
Copying data from file:/home/hadoop/Training/play_area/data/user-posts.txt
Copying file: file:/home/hadoop/Training/play_area/data/user-posts.txt
Loading data to table default.posts
Deleted /user/hive/warehouse/posts
OK
Time taken: 5.818 seconds
hive>
```

Existing records the table *posts* are deleted; data in *user-posts.txt* is loaded into Hive's *posts* table

```
$ hdfs dfs -cat /user/hive/warehouse/posts/user-posts.txt
user1,Funny Story,1343182026191
user2,Cool Deal,1343182133839
user4,Interesting Post,1343182154633
user5,Yet Another Blog,13431839394
```

Under the covers Hive stores it's tables in */user/hive/warehouse* (unless configured differently)

20

## 3: Query Data

```
hive> select count (1) from posts; ← Count number of records in posts table
Total MapReduce jobs = 1 ← Transformed HiveQL into 1 MapReduce Job
Launching Job 1 out of 1
```

```
...
Starting Job = job_1343957512459_0004, Tracking URL =
http://localhost:8088/proxy/application_1343957512459_0004/
Kill Command = hadoop job -Dmapred.job.tracker=localhost:10040 -kill
job_1343957512459_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2012-08-02 22:37:24,962 Stage-1 map = 0%, reduce = 0%
2012-08-02 22:37:30,497 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.87 sec
2012-08-02 22:37:31,577 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.87 sec
2012-08-02 22:37:32,664 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.64 sec
MapReduce Total cumulative CPU time: 2 seconds 640 msec
Ended Job = job_1343957512459_0004
MapReduce Jobs Launched:
Job 0: Map: 1 Reduce: 1 Accumulative CPU: 2.64 sec HDFS Read: 0 HDFS Write: 0
SUCESS
Total MapReduce CPU Time Spent: 2 seconds 640 msec
OK
4 ← Result is 4 records
Time taken: 14.204 seconds
```

21

## 3: Query Data

```
hive> select * from posts where user="user2";
```

```
...  
...
```

```
OK
```

```
user2 Cool Deal 1343182133839
```

```
Time taken: 12.184 seconds
```

Select records for "user2"

Select records whose  
timestamp is less or equals  
to the provided value

```
hive> select * from posts where time<=1343182133839 limit 2;
```

```
...  
...
```

```
OK
```

```
user1 Funny Story 1343182026191
```

```
user2 Cool Deal 1343182133839
```

```
Time taken: 12.003 seconds
```

```
hive>
```

Usually there are too  
many results to display,  
then one could utilize  
`limit` command to  
bound the display

22

## 4: Drop the Table

```
hive> DROP TABLE posts; ← Remove the table; use with caution
```

```
OK
```

```
Time taken: 2.182 seconds
```

```
hive> exit;
```

```
$ hdfs dfs -ls /user/hive/warehouse/
```

```
$
```

If hive was managing underlying file then it  
will be removed

23

# Loading Data

- **Several options to start using data in HIVE**

- Load data from HDFS location

```
hive> LOAD DATA INPATH '/training/hive/user-posts.txt'
> OVERWRITE INTO TABLE posts;
```

- File is copied from the provided location to /user/hive/warehouse/ (or configured location)

- Load data from a local file system

```
hive> LOAD DATA LOCAL INPATH 'data/user-posts.txt'
> OVERWRITE INTO TABLE posts;
```

- File is copied from the provided location to /user/hive/warehouse/ (or configured location)

- Utilize an existing location on HDFS

- Just point to an existing location when creating a table

24

# Re-Use Existing HDFS Location

```
hive> CREATE EXTERNAL TABLE posts
> (user STRING, post STRING, time BIGINT)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE
> LOCATION '/training/hive/';
```

OK

Time taken: 0.077 seconds

hive>

Hive will load all the files under  
/training/hive directory in posts table

25

# Schema Violations

- What would happen if we try to insert data that does not comply with the pre-defined schema?

```
hive> !cat data/user-posts-inconsistentFormat.txt;  
user1,Funny Story,1343182026191  
user2,Cool Deal,2012-01-05  
user4,Interesting Post,1343182154633  
user5,Yet Another Blog,13431839394
```

```
hive> describe posts;  
OK  
user  string  
post  string  
time  bigint  
Time taken: 0.289 seconds
```

Third Column 'post' is of type bigint;  
will not be able to convert  
'2012-01-05' value

26

# Schema Violations

```
hive> LOAD DATA LOCAL INPATH  
      > 'data/user-posts-inconsistentFormat.txt'  
      > OVERWRITE INTO TABLE posts;
```

```
OK  
Time taken: 0.612 seconds
```

```
hive> select * from posts;  
OK
```

```
user1 Funny Story 1343182026191  
user2 Cool Deal   NULL  
user4 Interesting Post 1343182154633  
user5 Yet Another Blog 13431839394  
Time taken: 0.136 seconds  
hive>
```

null is set for any value that  
violates pre-defined schema

27

# Partitions

- **To increase performance Hive has the capability to partition data**
  - The values of partitioned column divide a table into segments
  - Entire partitions can be ignored at query time
  - Similar to relational databases' indexes but not as granular
- **Partitions have to be properly created by users**
  - When inserting data must specify a partition
- **At query time, whenever appropriate, Hive will automatically filter out partitions**

28

## Creating Partitioned Table

```
hive> CREATE TABLE posts (user STRING, post STRING, time BIGINT)
> PARTITIONED BY(country STRING)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE;
```

Partition table based on the value of a country.

```
OK
Time taken: 0.116 seconds
```

```
hive> describe posts;
OK
```

```
user      string
post      string
time      bigint
country   string
```

There is no difference in schema between "partition" columns and "data" columns

```
Time taken: 0.111 seconds
```

```
hive> show partitions posts;
```

```
OK
Time taken: 0.102 seconds
hive>
```

29



# Load Data Into Partitioned Table

```
hive> LOAD DATA LOCAL INPATH 'data/user-posts-US.txt'
> OVERWRITE INTO TABLE posts;
FAILED: Error in semantic analysis: Need to specify partition
columns because the destination table is partitioned
```

Since the posts table was defined to be partitioned  
any insert statement must specify the partition

```
hive> LOAD DATA LOCAL INPATH 'data/user-posts-US.txt'
> OVERWRITE INTO TABLE posts PARTITION(country='US');
OK
Time taken: 0.225 seconds

hive> LOAD DATA LOCAL INPATH 'data/user-posts-AUSTRALIA.txt'
> OVERWRITE INTO TABLE posts PARTITION(country='AUSTRALIA');
OK
Time taken: 0.236 seconds
hive>
```

Each file is loaded into separate partition;  
data is separated by country

30

## Partitioned Table

- Partitions are physically stored under separate directories

```
hive> show partitions posts;
OK
country=AUSTRALIA
country=US
Time taken: 0.095 seconds
hive> exit;
```

There is a directory for  
each partition value

```
$ hdfs dfs -ls -R /user/hive/warehouse/posts
/user/hive/warehouse/posts/country=AUSTRALIA
/user/hive/warehouse/posts/country=AUSTRALIA/user-posts-AUSTRALIA.txt
/user/hive/warehouse/posts/country=US
/user/hive/warehouse/posts/country=US/user-posts-US.txt
```

31



## Querying Partitioned Table

- There is no difference in syntax
- When partitioned column is specified in the where clause entire directories/partitions could be ignored

Only "COUNTRY=US" partition will be queried,  
"COUNTRY=AUSTRALIA" partition will be ignored

```
hive> select * from posts where country='US' limit 10;
OK
user1 Funny Story 1343182026191      US
user2 Cool Deal  1343182133839      US
user2 Great Interesting Note 13431821339485  US
user4 Interesting Post 1343182154633      US
user1 Humor is good 1343182039586      US
user2 Hi I am user #2 1343182133839      US
Time taken: 0.197 seconds
```

32

## Bucketing

- Mechanism to query and examine random samples of data
- Break data into a set of buckets based on a hash function of a "bucket column"
  - Capability to execute queries on a sub-set of random data
- Doesn't automatically enforce bucketing
  - User is required to specify the number of buckets by setting # of reducer

```
hive> mapred.reduce.tasks = 256;
OR
hive> hive.enforce.bucketing = true;
```

Either manually set the # of reducers to be the number of buckets or you can use 'hive.enforce.bucketing' which will set it on your behalf

33

# Create and Use Table with Buckets

```
hive> CREATE TABLE post_count (user STRING, count INT)
      > CLUSTERED BY (user) INTO 5 BUCKETS;
OK
Time taken: 0.076 seconds
```

Declare table with 5 buckets for user column

```
hive> set hive.enforce.bucketing = true;
hive> insert overwrite table post_count
      > select user, count(post) from posts group by user;
Total MapReduce jobs = 2
Launching Job 1 out of 2
...
Launching Job 2 out of 2
...
OK
Time taken: 42.304 seconds
hive> exit;
$ hdfs dfs -ls -R /user/hive/warehouse/post_count/
/user/hive/warehouse/post_count/000000_0
/user/hive/warehouse/post_count/000001_0
/user/hive/warehouse/post_count/000002_0
/user/hive/warehouse/post_count/000003_0
/user/hive/warehouse/post_count/000004_0
```

# of reducer will get set 5

Insert data into post\_count bucketed table; number of posts are counted up for each user

A file per bucket is created; now only a sub-set of buckets can be sampled

34

# Random Sample of Bucketed Table


```
hive> select * from post_count TABLESAMPLE(BUCKET 1 OUT OF 2);
OK
user5  1
user1  2
Time taken: 11.758 seconds
hive>
```

Sample approximately 1 for every 2 buckets

35

# Joins

- **Joins in Hive are trivial**
- **Supports outer joins**
  - left, right and full joins
- **Can join multiple tables**
- **Default Join is Inner Join**
  - Rows are joined where the keys match
  - Rows that do not have matches are not included in the result



36

## Simple Inner Join

- **Let's say we have 2 tables: posts and likes**

```
hive> select * from posts limit 10;
```

OK

user1	Funny Story	1343182026191
user2	Cool Deal	1343182133839
user4	Interesting Post	1343182154633
user5	Yet Another Blog	1343183939434

Time taken: 0.108 seconds

```
hive> select * from likes limit 10;
```

OK

user1	12	1343182026191
user2	7	1343182139394
user3	0	1343182154633
user4	50	1343182147364

Time taken: 0.103 seconds

```
hive> CREATE TABLE posts_likes (user STRING, post STRING, likes_count INT);
```

OK

Time taken: 0.06 seconds

We want to join these 2 data-sets and produce a single table that contains user, post and count of likes

37

# Simple Inner Join

```
hive> INSERT OVERWRITE TABLE posts_likes  
> SELECT p.user, p.post, l.count  
> FROM posts p JOIN likes l ON (p.user = l.user);
```

OK

Time taken: 17.901 seconds

Two tables are joined based on user column; 3 columns are selected and stored in posts\_likes table

```
hive> select * from posts_likes limit 10;
```

OK

```
user1 Funny Story      12  
user2 Cool Deal        7  
user4 Interesting Post  50
```


Time taken: 0.082 seconds

hive>

38


# Outer Join

- Rows which will not join with the 'other' table are still included in the result



## Left Outer

- Row from the first table are included whether they have a match or not. Columns from the unmatched (second) table are set to null.



## Right Outer

- The opposite of Left Outer Join: Rows from the second table are included no matter what. Columns from the unmatched (first) table are set to null.



## Full Outer

- Rows from both sides are included. For unmatched rows the columns from the 'other' table are set to null.

39

# Outer Join Examples

```
SELECT p.*, l.*  
FROM posts p LEFT OUTER JOIN likes l ON (p.user = l.user)  
limit 10;
```

```
SELECT p.*, l.*  
FROM posts p RIGHT OUTER JOIN likes l ON (p.user = l.user)  
limit 10;
```

```
SELECT p.*, l.*  
FROM posts p FULL OUTER JOIN likes l ON (p.user = l.user)  
limit 10;
```

40

# Resources

- <http://hive.apache.org/>
- **Hive Wiki**
  - <https://cwiki.apache.org/confluence/display/Hive/Home>




## Hive

Edward Capriolo (Author), Dean Wampler (Author), Jason Rutherglen (Author)  
O'Reilly Media; 1 edition (October 3, 2012)


## Chapter About Hive Hadoop in Action

Chuck Lam (Author)  
Manning Publications; 1st Edition (December, 2010)



41

# Resources



**Chapter about Hive**  
**Hadoop in Practice**  
Alex Holmes (Author)  
Manning Publications; (October 10, 2012)



## Wrap-Up

**Customized Java EE Training:** <http://courses.coreservlets.com/>  
Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.  
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.



# Summary

- **We learned about**
  - Hive Concepts
  - Hive Installation
  - Table Creation and Deletion
  - Loading Data into Hive
  - Partitioning
  - Bucketing
  - Joins

44

© 2012 [coreservlets.com](http://coreservlets.com) and [Dima May](#)



## Questions?

[JSF 2, PrimeFaces, Java 7, Ajax, jQuery, Hadoop, RESTful Web Services, Android, Spring, Hibernate, Servlets, JSP, GWT, and other Java EE training.](#)

**Customized Java EE Training: <http://courses.coreservlets.com/>**  
Hadoop, Java, JSF 2, PrimeFaces, Servlets, JSP, Ajax, jQuery, Spring, Hibernate, RESTful Web Services, Android.  
Developed and taught by well-known author and developer. At public venues or onsite at *your* location.